

Floating-point numbers

Ágnes Baran, Csaba Noszály

Floating point numbers

Example 1

$a = 10$

$$0.3721 = \frac{3}{10} + \frac{7}{10^2} + \frac{2}{10^3} + \frac{1}{10^4}$$

$$21.65 = 0.2165 \cdot 10^2 = \left(\frac{2}{10} + \frac{1}{10^2} + \frac{6}{10^3} + \frac{5}{10^4} \right) \cdot 10^2$$

$a = 2$

$$0.1101 = \frac{1}{2} + \frac{1}{2^2} + \frac{0}{2^3} + \frac{1}{2^4}$$

$$0.001011 = 0.1011 \cdot 2^{-2} = \left(\frac{1}{2} + \frac{0}{2^2} + \frac{1}{2^3} + \frac{1}{2^4} \right) \cdot 2^{-2}$$

Floating point numbers

The form of non-zero floating point numbers:

$$\pm a^k \left(\frac{m_1}{a} + \frac{m_2}{a^2} + \cdots + \frac{m_t}{a^t} \right)$$

or in a shorter notation

$$\pm a^k \cdot 0.m_1 \dots m_t$$

or if a is known

$$\pm |k| m_1, \dots, m_t$$

where

$a > 1$ is an integer, the base,

$t > 1$ is an integer, the length of the mantissa

$k_- \leq k \leq k_+$ are integers, k is the characteristic, $k_- < 0$ and $k_+ > 0$ are fixed.

$0 \leq m_i \leq a - 1$ is an integer, for $i = 1, \dots, t$

If $m_1 > 0$ then the number is in *normalized* form. It makes the representation *unique*. Usually we will consider only normalized floating point numbers. The number 0 is not a normalized floating point number!

The set of the representable numbers is uniquely determined by the numbers

$$a, t, k_-, k_+$$

This set is denoted by $\mathcal{F}_{a,k_-,k_+,t}$ or simply by \mathcal{F} if the parameters are fixed.

Example 2

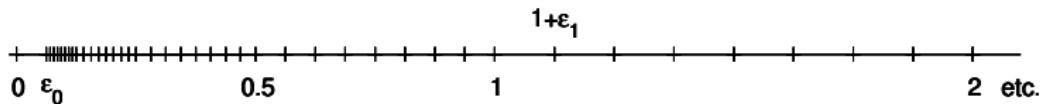
Let $a = 2$, $t = 4$, $k_- = -3$, $k_+ = 2$.

- Compute the floating-point form the numbers below:

$$0.6875, \quad 0.8125, \quad 3.25, \quad 0.875$$

- Of how many positive, normalized numbers can be represented in the given system?

For $a = 2$, $t = 4$, $k_- = -3$, $k_+ = 3$ the (partial) pictorial representation of the corresponding system is:



$$\epsilon_0 = a^{k_- - 1} = 2^{-4} = \frac{1}{16},$$

$$\epsilon_1 = a^{1-t} = 2^{-3} = \frac{1}{8}$$

Fact 1

For a given a, t, k_-, k_+

- the **largest** (positive) representable number:

$$\begin{aligned} M_\infty &= a^{k_+} \left(\frac{a-1}{a} + \frac{a-1}{a^2} + \cdots + \frac{a-1}{a^t} \right) = \\ &= a^{k_+} \left(1 - \frac{1}{a} + \frac{1}{a} - \frac{1}{a^2} + \cdots + \frac{1}{a^{t-1}} - \frac{1}{a^t} \right) = \\ &= a^{k_+} (1 - a^{-t}) \end{aligned}$$

- the **smallest** (positive) representable number:

$$\varepsilon_0 = a^{k_-} \left(\frac{1}{a} + 0 + \cdots + 0 \right) = a^{k_- - 1}$$

- **subnormal** numbers: if $k = k_-$ and $m_1 = 0$.

Fact 2

The number 1 is always representable:

$$1 = a^1 \cdot \frac{1}{a}$$

or

$$1 = [+|1|1, 0, \dots, 0]$$

The right neighbour of 1 (denoted by 1_+):

$$1 + \varepsilon_1 = [+|1|1, 0, \dots, 0, 1]$$

or

$$1 + \varepsilon_1 = a \left(\frac{1}{a} + 0 + \dots + 0 + \frac{1}{a^t} \right) = 1 + a^{1-t}$$

that is $\varepsilon_1 = a^{1-t}$ (**the machine epsilon**)

The **left** neighbour of 1 (denoted by 1_-):

$$1_- = a^0 \cdot 0.(a-1)\dots(a-1) = 1 - a^t$$

That is, the largest number from the previous characteristic. For powers of **a** the left and right neighbours are in different distance!

The IEEE floating point standard:

| | single precision | double precision |
|-----------------|------------------------------|-------------------------------|
| size | 32 bits | 64 bits |
| mantissa | 23+1 bits | 52+1 bits |
| characteristic | 8 bits | 11 bits |
| ε_1 | $\approx 1.19 \cdot 10^{-7}$ | $\approx 2.22 \cdot 10^{-16}$ |
| M_∞ | $\approx 10^{38}$ | $\approx 10^{308}$ |

Note that here m_1 is 1 (a constant), so it is not stored explicitly. For the sign 1 bit is reserved.

Example 3

The set of all positive normalized numbers in the system

$$a = 2, t = 4, k_- = -3, k_+ = 2$$

| | $k = -3$ | $k = -2$ | $k = -1$ | $k = 0$ | $k = 1$ | $k = 2$ |
|--------|------------------|-----------------|-----------------|-----------------|----------------|----------------|
| 0.1000 | $\frac{8}{128}$ | $\frac{8}{64}$ | $\frac{8}{32}$ | $\frac{8}{16}$ | $\frac{8}{8}$ | $\frac{8}{4}$ |
| 0.1001 | $\frac{9}{128}$ | $\frac{9}{64}$ | $\frac{9}{32}$ | $\frac{9}{16}$ | $\frac{9}{8}$ | $\frac{9}{4}$ |
| 0.1010 | $\frac{10}{128}$ | $\frac{10}{64}$ | $\frac{10}{32}$ | $\frac{10}{16}$ | $\frac{10}{8}$ | $\frac{10}{4}$ |
| 0.1011 | $\frac{11}{128}$ | $\frac{11}{64}$ | $\frac{11}{32}$ | $\frac{11}{16}$ | $\frac{11}{8}$ | $\frac{11}{4}$ |
| 0.1100 | $\frac{12}{128}$ | $\frac{12}{64}$ | $\frac{12}{32}$ | $\frac{12}{16}$ | $\frac{12}{8}$ | $\frac{12}{4}$ |
| 0.1101 | $\frac{13}{128}$ | $\frac{13}{64}$ | $\frac{13}{32}$ | $\frac{13}{16}$ | $\frac{13}{8}$ | $\frac{13}{4}$ |
| 0.1110 | $\frac{14}{128}$ | $\frac{14}{64}$ | $\frac{14}{32}$ | $\frac{14}{16}$ | $\frac{14}{8}$ | $\frac{14}{4}$ |
| 0.1111 | $\frac{15}{128}$ | $\frac{15}{64}$ | $\frac{15}{32}$ | $\frac{15}{16}$ | $\frac{15}{8}$ | $\frac{15}{4}$ |

$$M_\infty = 2^2(1 - 2^{-4}) = \frac{15}{4} \text{ and } \varepsilon_0 = 2^{-3-1} = \frac{1}{16} \left(= \frac{8}{128}\right)$$

Let $y = a^k \cdot 0.m_1m_2\dots m_t$.

The closest number that is greater than y is denoted by y_+ and:

$$y_+ = y + a^k \cdot \frac{1}{a^t} = y + a^{k-t}$$

This number is called the *stepsize* of the given characteristic.

Larger characteristic (exponent) means larger distance (stepsize) between neighbouring numbers.

If $k > t$, then the stepsize is larger than 1.

For double precision ($t = 53$):

| y | distance of the right neighbour |
|-------------------------------------|---------------------------------|
| 1 | $\approx 2.22 \cdot 10^{-16}$ |
| 16 | $\approx 3.5527 \cdot 10^{-15}$ |
| 1024 | $\approx 2.27 \cdot 10^{-13}$ |
| $2^{20} \approx 10^6$ | $\approx 2.33 \cdot 10^{-10}$ |
| $2^{52} \approx 4.5 \cdot 10^{15}$ | 1 |
| $2^{60} \approx 1.15 \cdot 10^{18}$ | 256 |
| $2^{66} \approx 7.38 \cdot 10^{19}$ | 16384 |

Rounding

Not all numbers has an exact representation in a floating point number system.

Example 4

The binary representation of $\frac{1}{10}$:

$$0.0001100110011001100\dots$$

The binary representation of $\frac{1}{3}$:

$$0.0101010101010\dots$$

Rounding

Let $x \in [-M_\infty, M_\infty]$ a real number, and denote by $fl(x)$ the corresponding floating-point number.

Regular rounding

$$fl(x) = \begin{cases} 0, & \text{if } |x| < \varepsilon_0 \\ \text{among the nearest floating point numbers to } x, \text{ the larger in absolute value,} & \text{if } |x| \geq \varepsilon_0 \end{cases}$$

Cutting, chopping

$$fl(x) = \begin{cases} 0, & \text{if } |x| < \varepsilon_0 \\ \text{the nearest floating point number towards zero,} & \text{if } |x| \geq \varepsilon_0 \end{cases}$$

Remark 1

The rounding rules implemented in todays processors are more involved.
For simplicity we will use the rules above.

Example 5

Let $a = 2$, $t = 4$, $k_- = -3$, $k_+ = 2$. What is $f\lceil(0.1)$ in case of chopping and regular rounding?

From the binary expansion of 0.1, we get the form:

$$2^{-3} \cdot 0.1100\textcolor{red}{110011001100\dots}$$

Regular rounding:

$$f\lceil(0.1) = 2^{-3} \cdot 0.1101$$

Chopping:

$$f\lceil(0.1) = 2^{-3} \cdot 0.1100$$

Rounding

Estimating the absolute error
in case of regular rounding:

$$|fl(x) - x| \leq \begin{cases} \varepsilon_0, & \text{ha } |x| < \varepsilon_0 \\ \frac{1}{2}\varepsilon_1|x|, & \text{ha } |x| \geq \varepsilon_0 \end{cases}$$

in case of chopping:

$$|fl(x) - x| \leq \begin{cases} \varepsilon_0, & \text{ha } |x| < \varepsilon_0 \\ \varepsilon_1|x|, & \text{ha } |x| \geq \varepsilon_0 \end{cases}$$

Rounding

Estimating the relative error
in case of regular rounding:

$$\frac{|fl(x) - x|}{|x|} \leq \frac{1}{2}\varepsilon_1$$

in case of chopping:

$$\frac{|fl(x) - x|}{|x|} \leq \varepsilon_1$$

Addition

Example 6

Let $a = 10$, $t = 3$. Assuming 1 spare digit compute $fl(x + y) = !$
 $x = 0.425 \cdot 10^{-1}$, $y = 0.677 \cdot 10^{-2}$

$$y \rightarrow y = 0.0677 \cdot 10^{-1} \quad (\mathbf{1 \ spare \ digit})$$

$$x + y = 0.425 \cdot 10^{-1} + 0.0677 \cdot 10^{-1} = 0.4927 \cdot 10^{-1}$$

$$fl(x + y) = \begin{cases} 0.492 \cdot 10^{-1}, & \text{chopping} \\ 0.493 \cdot 10^{-1}, & \text{regular rounding} \end{cases}$$

Error and operations

Denote by Δ one of the $+,-,*,/$, let x and y floating point numbers.
Assuming that the computer performs the operations exactly and assigns a floating point number to the result. Then
in case of regular rounding we have:

$$|fl(x\Delta y) - x\Delta y| \leq \begin{cases} \varepsilon_0, & \text{if } |x\Delta y| < \varepsilon_0 \\ \frac{1}{2}\varepsilon_1|x\Delta y|, & \text{if } |x\Delta y| \geq \varepsilon_0 \end{cases}$$

in case of chopping we have:

$$|fl(x\Delta y) - x\Delta y| \leq \begin{cases} \varepsilon_0, & \text{if } |x\Delta y| < \varepsilon_0 \\ \varepsilon_1|x\Delta y|, & \text{if } |x\Delta y| \geq \varepsilon_0 \end{cases}$$

$$\begin{aligned} |x \triangle y| > M_\infty &\implies \text{overflow} \\ |x \triangle y| < \varepsilon_0 &\implies \text{underflow} (fl(x \triangle y) = 0) \end{aligned}$$